




Early accuracy gains do not guarantee a joint qualification benefit

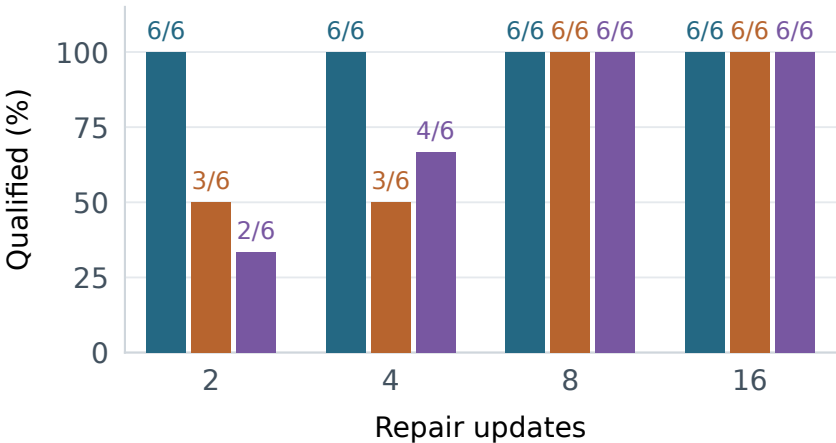
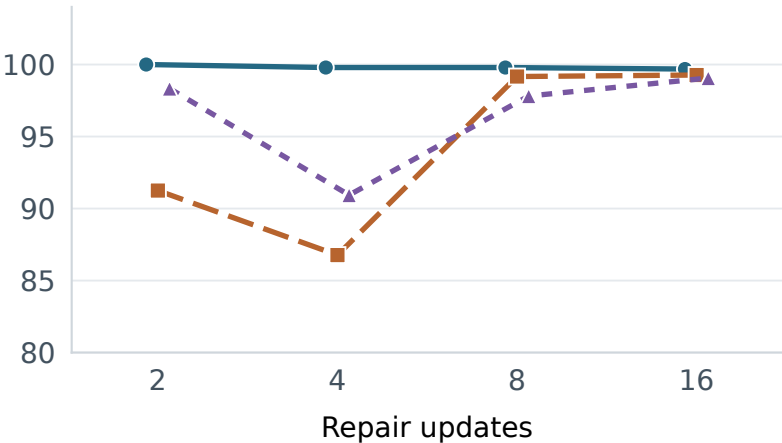
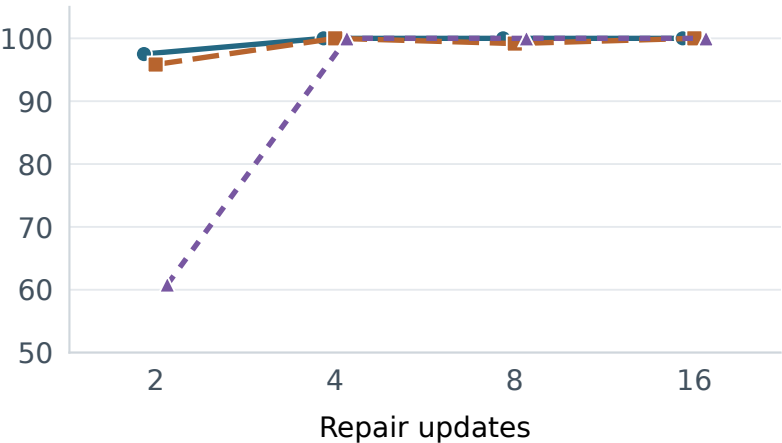
Held-out TEST sources shown separately · Consistent 16-code intent classification throughout

Original history
128 original + 128 new updates New-only
128 new updates; lifetime unmatched Alternate binding
128 rotated + 128 new; lifetime matched

TEST-A
7 paired intents
1 source model



TEST-B
6 paired intents
1 source model



At 8 updates, every history qualifies on all 13 intents: 7/7 in TEST-a and 6/6 in TEST-b.

Joint gate per intent: original-code accuracy $\geq 80\%$, gain ≥ 30 percentage points, and new-skill guard loss ≤ 10 points from its own baseline.

Protocol: same repair inputs/order; fresh AdamW from each history’s post-stream checkpoint per budget; 4 original + 4 new examples per update.

Two source-model clusters; the 13 intents are not independent replicates. Descriptive comparisons; no significance tests or confidence intervals.

Independent raw audit complete · 86,208 records (including repeated panels) · analysis 0f3477af996a... · audit d4c70e2dc036...