



Internal readouts did not improve repair allocation

Forecasts were sealed before held-out repairs. Every fixed budget was specified in advance.

Skills recovered while preserving new tasks



Updates selected by each policy



Two independent held-out models; 16-code intent classification. TRAIN-majority, output and confidence policies equal Fixed 2.

Original forecasts shown. One TRAIN label required an impossible gain; excluding it post hoc still gave no tuned-policy advantage.

All repair budgets were actually run. The right panel estimates policy allocations; it does not report saved study compute.